

Exact Solutions to NP-Complete Problems

Ref: - “Computer Algorithms”, Horowitz, Sahni, Rajasekaran (Chapters 7, 8)
- Various texts on Combinatorial Algorithms or on Integer Linear Programming

Backtracking

- An organized exhaustive search which often avoids searching many possibilities
- The desired solution is often expressed as n -tuple, where the x_i 's are chosen from some finite set S_i with $m_i = |S_i|$.
- The problem often requires finding one vector which maximizes, minimizes or satisfies a *criterion function* $P(x_1, x_2, \dots, x_n)$.
- The *brute force* approach is to evaluate each of $m = m_1 m_2 \dots m_n$ n -tuples from $S_1 \times S_2 \times \dots \times S_n$ and identify the n -tuple yielding the optimal value.
- The basis idea of backtracking is to build the solution vector using *modified criterion function* $P_i(x_1, x_2, \dots, x_i)$ to test whether the vector being formed have any chance of success.
- If a partial vector (a_1, a_2, \dots, a_i) has no chance of success, we avoid considering all of the $m_{i+1} m_{i+2} \dots m_n$ possible test vectors $(a_1, a_2, \dots, a_i, x_{i+1}, \dots, x_n)$
- Many problems solved by backtracking satisfy a set of *constraints* which may be divided into two categories: *explicit* and *implicit*.

Advanced Algorithms, Feodor F. Dragan, Kent State University

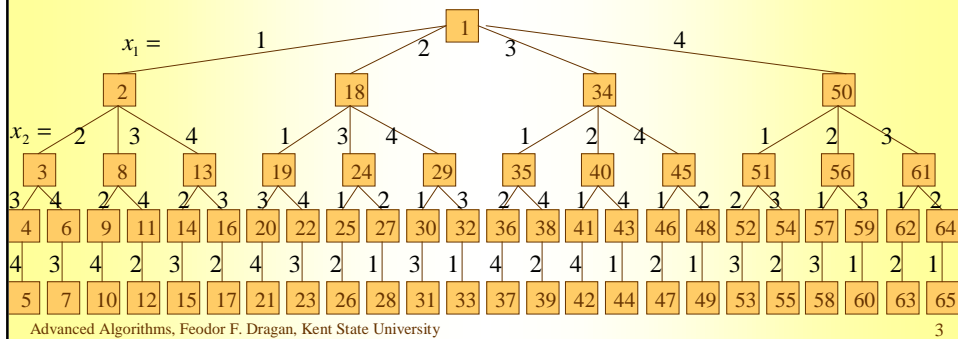
1

- **Explicit constraints** are rules that restrict each x_i to take on values only from a given set
 - examples: $x_i \geq 0$, $x_i = 0$ or 1 , $l_i \leq x_i \leq u_i$
 - depend on the particular instance I of the problem being solved
 - the n -tuples that satisfy these conditions define a possible *solution space* for I .
- **Implicit constraints** are rules which the tuples in the solution space for I must satisfy in order to satisfy the criterion function.
- **Example (8 Queens Problem)**
 - A classic problem in combinatorics is to place 8 queens on an 8 by 8 chessboard so that no two can “attack” each other (along a row, column, or diagonal).
 - Since each queen (1-8) must be on a different row, we can assume queen i is on row i .
 - All solutions to the 8-queens problem can be represented as an 8-tuple (x_1, x_2, \dots, x_8) where queen i is on column x_i .
 - The explicit constraints are $S_i = \{1, 2, \dots, 8\}$, $1 \leq i \leq 8$. The solution space consists of 8^8 8-tuples.
 - The implicit constraints are that no two x_i 's can be the same (as queens must be on different columns) and no two queens can be on the same diagonal.
 - this implies that all solutions are permutations of the 8-tuple $(1, 2, \dots, 8)$, and reduces the solution space from 8^8 tuples to $8!$ tuples.


Advanced Algorithms, Feodor F. Dragan, Kent State University

2

- Backtracking algorithms determine problem solutions by systematically searching the solution space.
- Search is facilitated using a *tree organization* for the solution space.
- Many tree organizations may be possible for the same solution space.
- **Example (*n*-Queens):** n queens are placed on n by n chessboard so that no two attack (no two queens are on the same row, column, or diagonal).
- Generalizing earlier discussion, solution space contains all $n!$ permutations of $(1,2,\dots,n)$.
- The tree below shows possible organization for $n=4$.
- Tree is called a *permutation tree* (nodes are numbered as in *depth first search*).
- Edges labeled by possible values of x_i
- The *solution space* is all paths from the root node to a leaf node
- There are $4!=24$ leaf nodes in tree



- **Example (*Sum of Subsets*):** Given positive numbers $w_i, 1 \leq i \leq n$, and m , find all subsets of $\{w_1, w_2, \dots, w_n\}$, whose sum is m .
- If $n=4$, $\{w_1, w_2, w_3, w_4\} = \{11, 13, 24, 7\}$ and $m=31$, the desired solution sets are $(11, 13, 7)$ and $(24, 7)$.
- If the solution vectors are given using the indices of the w_i values used, then the solution vectors are $(1, 2, 4)$ and $(3, 4)$.
- In general, all solutions are k -tuples (x_1, x_2, \dots, x_k) with $1 \leq k \leq n$ and different solutions may have different values of k .
- The *explicit constraints* on the solution space are that each $x_i \in \{1, 2, \dots, n\}$.
- The *implicit constraints* are that $x_i < x_{i+1}, 1 \leq i < n$, (so each item will occur only once) and that the sum of the corresponding w_i 's be m .




- The next figure gives the tree that corresponds to this variable tuple formation.
- An edge from a level i node to a level $i+1$ node represents a value for x_i .
- The solution space is all paths from the root node to any node in the tree.
- Possible paths include *empty path*, (1) , $(1, 2)$, $(1, 2, 3)$, $(1, 2, 3, 4)$, $(1, 2, 4)$, $(1, 3, 4)$, ...
- The leftmost subtree gives all subsets containing w_1 , the next subtree gives all subsets containing w_2 but not w_1 , etc

Advanced Algorithms, Feodor F. Dragan, Kent State University

4

- **Example (Sum of Subsets) again:** Another formulation of this problem represents each solution by an n -tuple (x_1, x_2, \dots, x_n) with $x_i \in \{0,1\}$, $1 \leq i \leq n$.
- Here $x_i = 0$ if w_i is not chosen and $x_i = 1$ if w_i is chosen.
- Given the earlier instance of $(11,13,24,7)$ and $m=31$, the solutions $(11,13,7)$ and $(24,7)$ are represented by $(1,1,0,1)$ and $(0,0,1,1)$.
- Here, all solutions have a fixed tuple size. The tree below corresponds to this formulation (nodes are numbered as in *D-search*).
- Edge from a level i node to a level $i+1$ node is labeled with the value of x_i (0 or 1)
- All paths from the root to a leaf give solution space.
- The left subtree gives all subsets containing w_1 and the right subtree gives all subsets not containing w_1 .



Advanced Algorithms, Feodor F. Dragan, Kent State University

5

Generating Problem States

- The two tree organizations for the sum of subsets problem are *static trees* (tree organization is independent of the problem instance being solved).
- Tree organizations that are problem instance dependent are called *dynamic trees* and are also used for some problems.
- Once a state space tree organization has been selected for a problem, the problem can be solved by
 - systematically generating the problem states,
 - finding which of these are solution states
 - finding which solution state are answer states
- A *live node* is a node which has been generated but whose children have not all been generated.
- An *E-node* (i.e., *expanding node*) is a live node whose children are currently being generated.
- A *dead node* is a generated node which is not to be expanded further or all of whose children have been generated.
- Two ways to generate problem states:
 - *Breadth First* Generation (queue of live nodes)
 - *Depth First* Generation (stack of live nodes)

Advanced Algorithms, Feodor F. Dragan, Kent State University

6

- **Depth First Generation** (stack of live nodes)
 - When a new child C of the current E-node R is generated, this child becomes the new E-node.
 - Then R will become the new E-node again when the subtree C has been fully explored.
 - Corresponds to a depth first search of the problem states.
- **Breadth First Generation** (queue of live nodes)
 - The E-node remains the E-node until it is dead.
- **Bounding functions** are used in both to kill live nodes without generating all of their children.
- At the end of the process, an answer node (or all answer nodes) are generated.
- The depth search generation method with bounding function is called *backtracking*.
- The breadth first generation method is used in the *branch-and-bound method*.

Example (backtracking on 4-queens problem)



- As a bounding function, use criterion that if (x_1, x_2, \dots, x_i) is the path to the current E-node, then some continuation $(x_1, x_2, \dots, x_i, x_{i+1}, \dots, x_n)$ exists that represents a chessboard without 2 queens attacking.
- Start with the root as the E-node. Then the path is $()$.
- The children of the E-nodes are generated in a left to right order.
- Node 2 is generated first and the path becomes (1). This corresponds to placing queen 1 on column 1.

Advanced Algorithms, Feodor F. Dragan, Kent State University

7

Example (backtracking on 4-queens problem)

- Node 2 becomes the E-node, as no two queens are attacking.
- Node 3 is generated and is immediately killed, as queens 1 and 2 would be on a diagonal.
- Node 8 is generated and the path (1,3) is ok, so node 8 becomes the next E-node.
- Node 8 gets killed since none of its children can lead to a feasible chessboard.
- Backtrack to node 2 and generate another child, node 13, giving path (1,4) which is ok.
- The first child of node 13 is node 14, which gives path (1,4,2) and the feasible chessboard.
- This process continues, as indicated in the figure. The figure shows the portion of the state space tree we had on slide #3.
- Note that only 16 out of 65 nodes (or 25%) in the solution space are actually generated.



Advanced Algorithms, Feodor F. Dragan, Kent State University

8

General Backtracking Algorithm

- This algorithm will find all answer nodes.
- If only the first solution is desired, a “flag” parameter can be added to indicate first success.
- Let (x_1, x_2, \dots, x_i) be a path from the root to a node in the state space tree.
- Let $T(x_1, x_2, \dots, x_i)$ be the set of all possible values for x_{i+1} such that $(x_1, x_2, \dots, x_i, x_{i+1})$ is also a path to a problem state (i.e., node).
- Let B_i be a boundary (Boolean) function such that if $B_i(x_1, x_2, \dots, x_i)$ is false, then the path (x_1, x_2, \dots, x_i) cannot be extended to reach an answer node.
- Note that if $B_i(x_1, x_2, \dots, x_i) = 1$, this does not guarantee that the path (x_1, x_2, \dots, x_i) can be extended to reach an answer node.
- Here is the recursive backtracking algorithm

```

Algorithm Backtrack(k)
for (each  $x[k]$  from  $T(x[1], x[2], \dots, x[k-1])$ ) do
{ if ( $B_k(x[1], x[2], \dots, x[k]) \neq 0$ ) then
  { if ( $x[1], x[2], \dots, x[k]$ ) is a path to an answer node)
    then write ( $x[1..k]$ );
    if ( $k < n$ ) then Backtrack( $k+1$ );
  }
}
    
```

Comments

- the candidates for $x[i+1]$ are values generated by $T(x[1], x[2], \dots, x[i])$ that satisfy B_{i+1}
- $T()$ gives all candidates for $x[1]$.
- elements are generated in a depth first manner, creating a preorder traversal (except for eliminated branches) of the state space tree.
- for many problems, the size of the state space tree is too large to permit generation of all nodes.

Advanced Algorithms, Feodor F. Dragan, Kent State University

9

Efficiency of Backtracking

- The efficiency of a backtracking algorithm depends upon 4 factors
 - the time to generate the next $x[k]$
 - the number of $x[k]$ choices that satisfy the explicit constraints
 - the time required to evaluate the bounding function B_i
 - the number of $x[k]$ satisfying the B_i
- A good boundary function will drastically reduce the number of candidates that have to be considered.
- Often a tradeoff between bounding functions, as one that is good may take more time to evaluate.
- For many problems such as n -queens, no good bounding function are known.
- Rearrangement:
 - the principle of selecting the set S_i with fewest elements each time
 - since these sets can be taken in any order, smaller branching at the higher levels create larger subtrees
 - removal of early nodes cut off larger subtrees (see Fig. 7.7 in HSR)
- The first three factors that effect the time required for backtracking depend primarily on the state space tree organization selected
- Only the fourth factor may vary widely, depending on the problem instances.

Advanced Algorithms, Feodor F. Dragan, Kent State University

10

- Worst case predictions for backtracking algorithms:
 - If the number of points in the solution space is 2^n or $n!$ the worst case timing is usually either $O(p(n)2^n)$ or $O(p(n)n!)$, where $p(n)$ is a polynomial.
 - Backtracking can often solve some problem instances with large n in very small amounts of time. However, may be difficult to predict behavior of algorithm for particular problem instances.
- Estimating Nr. of nodes generated:
 - The number of nodes generated in a particular instance can be estimated using Monte Carlo methods
 - Starting at the top level, a random path is generated, as follows:
 - set x be a node on this path at level i of the state space tree; the boundary function B_i is used to determine the number m_i of its children which will be generated; one child is randomly selected, and the process continues until the path ends.
 - Then $m = m_1 + m_1m_2 + m_1m_2m_3 + \dots$ is an estimate of the nodes that will be generated.
 - Above estimate for m assumes the bounding functions are static and do not improve with time; It also assumes that the same boundary function is used for all nodes at the same level.
 - The above two assumptions are not true for most backtracking algorithms; e.g., the boundary functions usually get stronger as information is gathered about the search.
 - Consequently, the above value of m is likely to be high when these two assumptions are false.
 - A better estimate would also result if the value m is the average returned for several (about 20) random paths.

Advanced Algorithms, Feodor F. Dragan, Kent State University

11

Backtracking Algorithm for n-Queens problem

- Let (x_1, x_2, \dots, x_n) represent where the i th queen is placed (in row i and column x_i) on an n by n chessboard.
- Observe that two queens on the same diagonal that runs from “upper left” to “lower right” have the same “row-column” value.
- Also two queens on the same diagonal from “upper-right” to “lower left” have the same “row+column” value
- Then two queens at (i, j) and (k, l) are on the same diagonal if and only if

$$i-j=k-l \text{ or } i+j=k+l$$
 iff

$$i-k=j-l \text{ or } j-l=k-i$$
 iff

$$|j-l|=|i-k|.$$
- Algorithm PLACE(k, i) returns *true* if the k th queen can be placed in column i and runs in $O(k)$ time (see next slide)
- Using PLACE, the recursive version of the general backtracking method can be used to give a precise solution to the n-queens problem
- Array $x[1..n]$ is global. Algorithm invoked by NQUEENS(1, n).

Advanced Algorithms, Feodor F. Dragan, Kent State University

12

```

bool Place(int k, int i)
// Returns true if a queen can be placed in kth row and ith column. Otherwise it returns false.
// x[] is a global array whose first (k-1) values have been set.
// abs(r) returns the absolute value of r.
{
    for (int j=1; j < k; j++)
        if ((x[j] == i) // Two in the same column
            || (abs(x[j]-i) == abs(j-k))) // or in the same diagonal
            return(false);
    return(true);
}

void NQueens(int k, int n)
// Using backtracking, this procedure prints all possible placements of n queens on an n x n
// chessboard so that they are nonattacking.
{
    for (int i=1; i<=n; i++) {
        if (Place(k, i)) {
            x[k] = i;
            if (k==n) { for (int j=1; j<=n; j++)
                        cout << x[j] << ' '; cout << endl; }
            else NQueens(k+1, n);
        }
    }
}

```

Advanced Algorithms, Feodor F. Dragan, Kent State University

13

Efficiency of n-Queens over Brute Force

- For an 8x8 chessboard, there are $\binom{64}{8}$ ways to place 8 queens on the chessboard (billions of 8-tuples to examine)
- Requiring placement of queens on distinct rows and columns reduces the number of 8-tuples that must be examined to $8! = 40,320$ 8-tuples
- Next we estimate the number of nodes that will be generated by NQUEENS. The assumptions needed for this estimate hold for NQUEENS.
 - Boundary function is static
 - boundary functions does not change as search progress
 - additionally, nodes on the same level of the tree have the same degree
- Five trials using the ESTIMATE function described earlier are given on p.355 of HSR
 - each produces a random path and estimates the total number of nodes generated, based on this path
 - the average of these 5 estimates is 1625
 - the total number of nodes in the 8-queens state space tree is

$$1 + \sum_{j=0}^7 [\prod_{i=0}^j (8-i)] = 69,281$$
 - the estimated number of unbound nodes is only 2.34 % of the total number of nodes in the 8-queens state space tree.

	1						
			2				
3							
		4					
				5			

(8,5,4,3,2)=1649

Advanced Algorithms, Feodor F. Dragan, Kent State University

14

Sum of Subsets Algorithm Overview

- **Problem Restated:** Given n distinct positive integers (called weights), find all combinations of these numbers whose sum is m .
- We use the state space tree based on the fixed tuple length (w_1, w_2, \dots, w_n) where $x_i = 0$ if w_i is not included and $x_i = 1$ if w_i is included (see Fig. on slide #5).
- The weights (w_1, w_2, \dots, w_n) are assumed to initially be sorted into increasing order.
- Note that the tree node corresponding to (x_1, x_2, \dots, x_k) cannot lead to an answer node unless

$$(1) \quad \sum_{i=1}^k w_i x_i + \sum_{i=k+1}^n w_i \geq m$$

- Also, note that (x_1, x_2, \dots, x_k) cannot lead to an answer node unless

$$(2) \quad \sum_{i=1}^k w_i x_i + w_{k+1} \leq m$$



- The boundary function used uses both of the preceding conditions:

$$B_k(x_1, x_2, \dots, x_k) = \text{true} \text{ iff } (1) \text{ and } (2) \text{ hold}$$

- The algorithm for the sum of subsets problem given in HSR is obtained by using this boundary function in the general recursive backtracking algorithm.
- A couple of implementation simplifications used are explained in HSR (see pp. 358-9).

Hamiltonian Cycles Algorithm

- Example: Graph $G1$ contains a Hamiltonian cycle 1,2,8,7,6,5,4,3,1 while graph $G2$ contains no Hamiltonian cycle.



- The algorithm given works on both directed and undirected graphs
- All distinct cycles will be found.
- Each x_i in the backtracking solution vector (x_1, x_2, \dots, x_n) represents the i th vertex visited in the proposed cycle
- The vertices of the graph are assumed to be named using the first n positive integers.
- To avoid printing the same cycle n times, we require $x_1 = 1$
- The algorithm *NextValue* determines a possible next vertex for the proposed cycle
 - if $1 < k < n$, then x_k can be any vertex v that is distinct from x_1, x_2, \dots, x_{k-1} and is connected by an edge to x_{k-1}
 - the vertex x_n must be one remaining vertex and must be connected by an edge to both x_1 and x_{n-1}
- The backtracking algorithm *Hamiltonian(k)* is obtained by using *NextValue* to select a legal vertex to add. No boundary function is used.

- The main algorithm starts by
 - initializing the adjacency matrix $G[1:n, 1:n]$
 - setting $x[2:n]=0$
 - setting $x[1]=1$
 - executing $\text{Hamiltonian}(2)$

Recursive
algorithm that
finds all
Hamiltonian
cycles



```
void Hamiltonian(int k)
// This program uses the recursive formulation of
// backtracking to find all the Hamiltonian cycles
// of a graph. The graph is stored as an adjacency
// matrix G[1:n][1:n]. All cycles begin at node 1.
{
    do { // Generate values for x[k].
        NextValue(k); // Assign a legal next value to x[k].
        if (!x[k]) return;
        if (k == n) {
            for (int i=1; i<=n; i++) cout << x[i] << ' ';
            cout << "\n";
        }
        else Hamiltonian(k+1);
    } while (1);
}
```

```
void NextValue(int k)
// x[1],...,x[k-1] is a path of k-1 distinct vertices.
// If x[k]==0, then no vertex has as yet
// been assigned to x[k]. After execution x[k] is assigned
// to the lowest numbered vertex which
// i) does not already appear in x[1],x[2],...,x[k-1]; and
// ii) is connected by an edge to x[k-1].
// Otherwise x[k]==0.
// If k==n, then in addition x[k] is connected to x[1].
{do {
    x[k] = (x[k]+1) % (n+1); // Next vertex
    if (!x[k]) return;
    if (G[x[k-1]][x[k]]) { // Is there an edge?
        for (int j=1; j<=k-1; j++) if (x[j]==x[k]) break;
        // Check for distinctness.
        if (j==k) // If true, then the vertex is distinct.
            if ((k<n) || ((k==n) && G[x[n]][x[1]]))
                return;
    }
} while(1);
}
```

Returning at this line
causes $\text{Hamiltonian}(k)$ to
backtrack

Assigns $x[k]$ the values
1,2,...,n successively,
following failure at

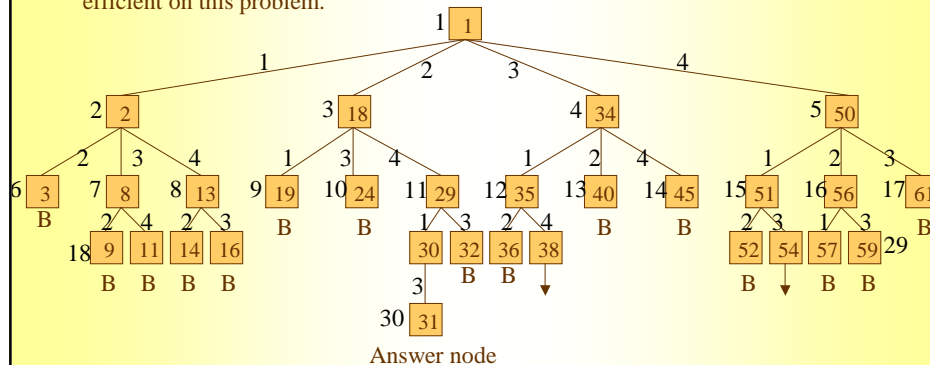
Homework problem:

Generalize Hamiltonian so
that it processes a graph
whose edges have cost
associated with them and
finds a Hamiltonian cycle
with minimum cost. You can
assume that all edge costs are
positive (Ex. 3, p. 368 of HSR).

Branch and Bound Algorithms

- Refers to all state space search methods in which all children of the E-nodes are generated before any other live nodes become E-nodes.
- **Breadth First Search** (BFS) will be called FIFO (first in, first out)
 - each new node is placed in a queue
 - after all children of current E-node are generated, the node at front of queue become the E-node.
- **D-search** will be called a LIFO (last in, first out)
 - new nodes are placed in stack
- As with backtracking, bounding functions will be used to avoid generating trees with no answer node.
- **Example:** 4 Queen FIFO Branch & Bound Algorithm.
 - The state space tree in Fig. on slide #3 is used, so node numbers do not indicate order of generation.
 - Initially, only the root node is alive (no queens placed)
 - Expanding the root E-node generates its children nodes in the order 2,18,34, and 50. These nodes represent a 4x4 chessboard with queen 1 in row 1 and columns 1,2,3, and 4 respectively.
 - The only live nodes are now 1,18,34,50, and the next E-node is 2. It is expanded, generating nodes 3,8,13.

- Node 3 is killed immediately by the bounding function used in backtracking algorithm and nodes 8,13 are added to queue of live nodes.
- This process is continued, generating below figure.
- Comparing the two trees of generated nodes, it is clear that backtracking is more efficient on this problem.



Least Cost (LC) search

- Both FIFO and LIFO are rigid and blind.
- The search for an answer node can often be speeded up using an “intelligent” ranking function $C()$ for the live nodes.

Least Cost (LC) search

- In 4-Queens example, if $C()$ had assigned node 30 a better rank than other live nodes, it would have become the E-node following node 29.
- An ideal way to assign rank to each live node x is the number of levels the nearest answer node (in the subtree with root x) is from x .
- Using this ranking function on the previous 4-queens example would have assigned rank 1 to both answer nodes 30 and 38.
- Let $g(x)$ be an estimate of the additional effort needed to reach an answer node from x .
- Node x is assigned a rank using a function $C()$ defined by $C(x) = f(h(x)) + g(x)$, where $h(x)$ is the cost of reaching x from the root.
 f is a non-decreasing weight function.
- The use of nonzero f helps prevent the search algorithm from making unnecessarily deep probes into the search tree.
- A nonzero f is needed, as otherwise a child y of the current E-node x will become the next E-node since $g(y) \leq g(x)$ and x had the previous lowest rank.
- Use of non-zero f forces the search algorithm to favor nodes closer to the root, reducing the probability of a deep and fruitless search into the tree.
- A search that uses a cost function $C()$ to choose the next E-node to be a live node with minimum $C()$ value is called a **LC-search**.

- If $g=0$, $f=1$, and $h(x)$ is the level of node x , this LC-search is a BFS algorithms which generates nodes by levels.
- If $f=0$ and $g(y) < g(x)$ when y is a child of x , this LC-search is a D-search.
- The *cost function* $c()$ is defined as follows
 - if x is an answer node, then $c(x)$ is the cost of reaching x from the root of the state space tree.
 - if x is not an answer node, but the subtree of x contains an answer node, then $c(x)$ is the minimal cost of an answer node in subtree x .
 - otherwise, $c(x) = \infty$
- Then, $C()$ with $f=1$, that is, $C(x) = h(x) + g(x)$, is an estimate of $c()$.
- $C()$ should be chosen so that it is easy to compute. It will normally have the additional property that if x is an answer node or leaf node, then $c(x) = C(x)$.

Least Cost Search Algorithm

- This algorithm assumes two additional algorithms, *Least(x)* and *Add(x)* to manage the list of live nodes.
- *Least()* finds a live node with least $C()$ value. This node is deleted from the list of live nodes and returned.
- *Add(x)* adds the new live node x to the list of live nodes.
- With each node x that becomes alive, we associate a field *parent* which stores the parent of x .

Least Cost Search Algorithm

- This allows LC-search to output a path from the answer node it finds to the root node.
- LC-search terminates only when either an answer node is found or the entire state space tree has been generated and searched.
- Note that termination is only guaranteed for finite space trees.
- It is advisable to restrict the search in LC-search to find answer nodes with costs not exceeding a given bound C .
- Note: *Least* and *Add* can be defined to implement a stack or queue as well, so the algorithms for LC, FIFO, and LIFO search are essentially the same.

```
struct listnode { struct listnode *next, *parent; float cost; };
LCSearch(struct listnode *t) // Search t for an answer node.
{
    Struct listnode *x, *E, *Least();
    if (*t is an answer node) output *t and return;
    E=t; // E-node
    initialize the list of live nodes to be empty;
    do { for (each child x of E) {
        if (x is an answer node) output the path from x to t and return;
        Add(x); // x is a new live node.
        x->parent = E; // pointer for path to root
    }
    if there are no more live nodes { count << "No answer node\n"; return; }
    E=Least();
} while(1);
}
```

Advanced Algorithms, Feodor F. Dragan, Kent State University

23

Dynamic Programming

- This is another technique for finding exact solutions to NP-Complete problems.
- Examples:
 - 0/1 Knapsack Problem (your old homework problem; $O(nW)$ time algorithm)
 - Traveling Salesman Problem (HSR p. 298)

READ (in HSR) Sections: 7.1-7.3, 7.5, and 8.1(8.1.1-8.1.3).

Homework 7:

• Problems:

- no more problems at this moment

Advanced Algorithms, Feodor F. Dragan, Kent State University

24